

Run Qwen 3.8 Flash Next (125B) on consumer hardware (RTX 4090) at 100T/s

From: Hacker News

Source: Hacker News

The Brief

****What shipped****

Qwen 3.8 Flash Next, a 125-billion-parameter AI model, is available for use on consumer hardware.

****Capabilities****

The model can chat, write code, read images, and interact with applications, with a context length of up to 32K tokens. Performance benchmarks show it can write answers at speeds of up to 94 tokens per second on an RTX 5070 and read prompts at speeds of up to 2,650 tokens per second.

****Availability****

Strata, the software to run the model, is free and open source, available for download from GitHub. It requires a compatible graphics card with at least 12 GB of VRAM and can be set up on Windows or Linux.

****Why it matters****

This allows developers to run a powerful AI model locally, reducing reliance on cloud services and potentially lowering costs associated with API usage. It also enables faster iterations and testing of AI applications directly on consumer-grade hardware.

****Caveats****

Performance may vary significantly based on hardware specifications, and the model's capabilities are dependent on the available RAM and VRAM. Some features, like reading images on AMD cards, are still experimental.

Full Announcement Details

Qwen 3.8 Flash Next, a 125-billion-parameter AI model, is available for use on consumer hardware. The model can chat, write code, read images, and interact with applications, with a context length of up to 32K tokens. Performance benchmarks show it can write

Daily AI Dev Digest

Sunday, October 04, 2026

answers at speeds of up to 94 tokens per second on an RTX 5070 and read prompts at speeds of up to 2,650 tokens per second.

Want to learn more?

<https://github.com/Niko1221/Strata>

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.