

My research agenda and work

By: AI Alignment Forum

Source: AI Alignment Forum

What's New

****Claim****

The research aims to predict the characteristics and failure modes of the first transformative AI to inform alignment strategies.

****Evidence****

The researcher has been engaged in AI alignment work for over three years and has been contemplating brainlike AGI and its alignment since 2004. This extensive background suggests a long-term investment in understanding the mechanisms of transformative AI and its implications for alignment.

****Method****

The approach involves mechanistic predictions about transformative AI, which likely includes theoretical modeling and analysis of potential alignment failure modes.

****Limitations****

The scope of the predictions may be limited by the current understanding of AI capabilities and alignment challenges, and there may be gaps in reproducibility due to the speculative nature of predicting future AI systems.

****Safety relevance****

This research updates the understanding of potential alignment challenges and informs methods of evaluation and governance for future AI systems.

Full Announcement Details

The research aims to predict the characteristics and failure modes of the first transformative AI to inform alignment strategies. The researcher has been engaged in AI alignment work for over three years and has been contemplating brainlike AGI and its alignment since 2004. This extensive background suggests a long-term investment in understanding the mechanisms of transformative AI and its implications for alignment.

Daily AI Dev Digest

Friday, June 05, 2026

Want to learn more?

<https://www.alignmentforum.org/posts/MuLvZxMcy5WaKJu3H/my-research-agenda-and-work>

This digest tracks the latest developer-facing product features from top AI labs. Stay ahead by trying new APIs, models, and tools as they launch.