

Daily AI Dev Digest

Monday, October 05, 2026

llama.cpp v0.6.0 released

From: *GitHub*

Source: *GitHub*

The Brief

****What shipped****

llama.cpp v0.6.0 was released, introducing new features and models including the GLM-5.3-Flash (GLM5-Next) and Clef decision model.

****Capabilities****

The new GLM-5.3-Flash is a 320B hybrid text+vision model, while the Clef decision model supports both text and vision inputs. The update includes a new llama_batch_ext API for mixed token/embedding inputs and MTP/deepstack state embeddings, as well as MTP speculative decoding for Qwen4Exp, which reportedly achieves a ~1.5x decode speedup on DGX Spark.

****Availability****

The updated llama.cpp v0.6.0 and its features are available on GitHub. Specific pricing and access details are not stated.

****Why it matters****

The introduction of the llama_batch_ext API allows for more flexible input handling, which can improve the efficiency of model serving and integration into applications. The performance improvements in decoding speed may also reduce latency in production environments.

****Caveats****

Claims regarding speed improvements and model performance have not been independently verified, and specific benchmark results are not provided. Additionally, access to the new models and features may depend on the user's setup and environment.

Full Announcement Details

llama.cpp v0.6.0 was released, introducing new features and models including the GLM-5.3-Flash (GLM5-Next) and Clef decision model. The new GLM-5.3-Flash is a 320B hybrid text+vision model, while the Clef decision model supports both text and vision

Daily AI Dev Digest

Monday, October 05, 2026

inputs. The update includes a new llama_batch_ext API for mixed token/embedding inputs and MTP/deepstack state embeddings, as well as MTP speculative decoding for Qwen4Exp, which reportedly achieves a ~1.5x decode speedup on DGX Spark.

Want to learn more?

<https://github.com/ggml-org/llama.cpp/releases/tag/v0.6.0>

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.