

EmbeddingGemma 2: an open, lightweight multimodal embedding model

From: Google DeepMind

Source: Google DeepMind

The Brief

****What shipped****

Google DeepMind shipped EmbeddingGemma 2, an open, lightweight multimodal embedding model.

****Capabilities****

EmbeddingGemma 2 can process text, images, audio, and video in a unified embedding space with 740 million parameters. It features an 8K token context window and achieves a 9.92-point improvement in code performance on the MTEB Code benchmark, scoring 78.68.

****Availability****

The model weights are available for download on Hugging Face and Kaggle, with Gemini Enterprise Agent Platform Model Garden availability coming soon. Not stated if there are any pricing tiers.

****Why it matters****

This model enables developers to build robust on-device multimodal applications that ensure data privacy and reduce latency, allowing for efficient local search and retrieval capabilities.

****Caveats****

Claims regarding performance improvements and benchmarks are unverified, and specific details on the model's limitations or potential issues are not provided.

Full Announcement Details

Google DeepMind shipped EmbeddingGemma 2, an open, lightweight multimodal embedding model. EmbeddingGemma 2 can process text, images, audio, and video in a unified embedding space with 740 million parameters. It features an 8K token context window and achieves a 9.92-point improvement in code performance on the MTEB Code benchmark, scoring 78.68.

Daily AI Dev Digest

Tuesday, October 06, 2026

Want to learn more?

<https://deepmind.google/blog/embeddinggemma-2-an-open-lightweight-multimodal-embedding-model/>

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.