Daily AI Dev Digest

Saturday, May 09, 2026

Running Codex safely at OpenAI

By: OpenAI

Source: OpenAI

What's New

Claim
OpenAI employs a range of security measures to operate Codex safely and ensure compliance in coding tasks.

Evidence
OpenAI's approach includes implementing sandboxing techniques, requiring approvals for certain actions, establishing network policies, and utilizing telemetry specific to the agent's behavior. These measures are designed to monitor and mitigate potential risks associated with Codex's deployment.

Method
The methodology involves a combination of technical controls, such as sandboxing and network restrictions, along with real-time telemetry to track the system's behavior and performance. A structured approval process is in place for actions that could pose higher risks.

Limitations
The description does not provide specific quantitative data on the effectiveness of these safety measures or their impact on real-world scenarios. Additionally, the scope appears limited to Codex without addressing broader implications for AI safety across different models.

Safety relevance
This approach informs the ongoing discussion on AI alignment and governance by highlighting practical safety measures that can be implemented to manage risks in coding agents.

Full Announcement Details

OpenAI employs a range of security measures to operate Codex safely and ensure compliance in coding tasks. OpenAI's approach includes implementing sandboxing techniques, requiring approvals for certain actions, establishing network policies, and utilizing telemetry specific to the agent's behavior. These measures are designed to monitor and mitigate potential risks associated with

Codex's deployment.

Want to learn more?

https://openai.com/index/running-codex-safely

This digest tracks the latest developer-facing product features from top AI labs. Stay ahead by trying new APIs, models, and tools as they launch.