

SFT Drives Gemini's Safety Properties

By: AI Alignment Forum

Source: AI Alignment Forum

What's New

****Claim****

Most safety-relevant properties in the Gemini model are primarily attributed to the combination of pretraining and supervised fine-tuning (SFT), rather than reinforcement learning (RL).

****Evidence****

The post indicates that the authors observed a surprising correlation between safety properties and the training stages, particularly noting that the safety aspects emerged more strongly from pretraining and SFT. Specific empirical data or statistical analyses are not detailed in the summary.

****Method****

The findings are based on an analysis of the training stages of the Gemini model, comparing the effects of pretraining, SFT, and RL on safety properties, although the exact experimental setup is not described.

****Limitations****

The authors caution against overgeneralizing these findings to other model families and acknowledge potential gaps in how comprehensively their results cover different training methodologies or model architectures.

****Safety relevance****

This update suggests that the interaction between pretraining and SFT may be more critical for safety in language models than previously understood, informing future alignment and evaluation strategies.

Full Announcement Details

Most safety-relevant properties in the Gemini model are primarily attributed to the combination of pretraining and supervised fine-tuning (SFT), rather than reinforcement learning (RL). The post indicates that the authors observed a surprising correlation between safety properties and the training stages, particularly noting that the safety aspects emerged more strongly from pretraining and SFT. Specific empirical data or statistical analyses are not detailed in the summary.

Daily AI Dev Digest

Saturday, June 13, 2026

Want to learn more?

<https://www.alignmentforum.org/posts/nLrrYweeFvgXACSmS/sft-drives-gemini-s-safety-properties-1>

This digest tracks the latest developer-facing product features from top AI labs. Stay ahead by trying new APIs, models, and tools as they launch.