Daily AI Dev Digest

Thursday, May 14, 2026

The safe-to-dangerous shift is a fundamental problem for eval realism; but also for measuring awareness

By: AI Alignment Forum

Source: AI Alignment Forum

What's New

Claim
The safe-to-dangerous shift poses a significant challenge for evaluating AI alignment realism.

Evidence
The discussion highlights that if a capable model can distinguish between deployment and evaluation distributions, alignment evaluations may not provide reliable assurance against catastrophic outcomes. This suggests a mismatch between evaluation conditions and real-world deployment scenarios.

Method
The work primarily involves theoretical analysis of black-box alignment evaluations and their limitations in distinguishing deployment contexts from evaluation contexts.

Limitations
The analysis may not be generalizable across all AI models and does not provide empirical data from specific experiments or datasets to validate the claims made.

Safety relevance
This underscores the need for more robust evaluation frameworks that account for potential shifts in model behavior upon deployment, impacting alignment strategies.

Full Announcement Details

The safe-to-dangerous shift poses a significant challenge for evaluating AI alignment realism. The discussion highlights that if a capable model can distinguish between deployment and evaluation distributions, alignment evaluations may not provide reliable assurance against catastrophic outcomes. This suggests a mismatch between evaluation conditions and real-world deployment scenarios.

Daily AI Dev Digest

Thursday, May 14, 2026

Want to learn more?

https://www.alignmentforum.org/posts/tK8vqHDxaRGcysNJQ/the-safe-to-dangerous-shift-is-a-fundamental-problem-for-1

This digest tracks the latest developer-facing product features from top AI labs. Stay ahead by trying new APIs, models, and tools as they launch.