Daily AI Dev Digest

Friday, October 16, 2026

Scaling Indic Pretraining: 2.8 Trillion Tokens and an Open Data Generation Pipeline

From: Hugging Face Blog

Source: Hugging Face Blog

The Brief

What shipped
AI4Bharat and NVIDIA shipped the Nemotron-pretraining-Indic-v1 model, which includes approximately 2.8 trillion tokens across 22 Indian languages.

Capabilities
The model supports native-script and romanized data for 22 Indian languages, with a total of 2,601.0 billion native-script tokens and 224.6 billion romanized tokens. It utilizes the Gemma-4 tokenizer and incorporates a four-stage synthetic data generation pipeline that includes translation, quality assessment, and targeted repair.

Availability
The model and the open GPU-accelerated data generation pipeline will be released soon, but specific details on access, such as API availability or pricing, are not stated.

Why it matters
This release provides a substantial resource for building models in Indic languages, addressing the data scarcity issue and enabling developers to generate more training data tailored to their needs.

Caveats
The announcement does not provide specific benchmark numbers or performance metrics for the model, and details on the exact process for accessing the resources are not stated.

Full Announcement Details

AI4Bharat and NVIDIA shipped the Nemotron-pretraining-Indic-v1 model, which includes approximately 2.8 trillion tokens across 22 Indian languages. The model supports native-script and romanized data for 22 Indian languages, with a total of 2,601.0 billion

native-script tokens and 224.6 billion romanized tokens. It utilizes the Gemma-4 tokenizer and incorporates a four-stage synthetic data generation pipeline that includes translation, quality assessment, and targeted repair.

Want to learn more?

https://huggingface.co/blog/nvidia/scaling-indic-pretraining

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.