

Safety and alignment in an era of long-horizon models

By: OpenAI

Source: OpenAI

What's New

****Claim****

Long-running AI models introduce new safety risks and observed failures that necessitate improved safeguards.

****Evidence****

OpenAI reported specific failures encountered during the deployment of long-running models, highlighting incidents that led to unintended behaviors. They also noted iterative improvements in safety measures based on these experiences, though quantifiable metrics were not provided.

****Method****

The work involved deploying AI models over extended periods and observing their performance in real-world scenarios to identify safety risks. Iterative deployment strategies were employed to refine safety protocols.

****Limitations****

The findings may not generalize across all types of AI models, and specific quantitative data supporting the claims were lacking, which raises concerns about reproducibility and the robustness of the conclusions.

****Safety relevance****

This work emphasizes the need to adapt safety frameworks to account for the unique challenges posed by long-horizon AI systems.

Full Announcement Details

Long-running AI models introduce new safety risks and observed failures that necessitate improved safeguards. OpenAI reported specific failures encountered during the deployment of long-running models, highlighting incidents that led to unintended behaviors. They also noted iterative improvements in safety measures based on these experiences, though quantifiable metrics were not provided.

Daily AI Dev Digest

Monday, July 20, 2026

Want to learn more?

<https://openai.com/index/safety-alignment-long-horizon-models>

This digest tracks the latest developer-facing product features from top AI labs. Stay ahead by trying new APIs, models, and tools as they launch.