Daily AI Dev Digest

Saturday, October 31, 2026

Le Chonk: Built for Coders, Trusted by Defenders.

Source: Mistral AI

The Brief

What shipped
Mistral AI released the open weights for its improved Mistral Large 4 (Le Chonk), a 1-trillion-parameter natively multimodal model.

Capabilities
Mistral reports a 19.7% score on Cognition's FrontierCode 2.0 and 48% on the Artificial Analysis Coding Agent Index, tied with GLM-5.3 while using 47% fewer tokens per task on average. It scores 41% on Terminal-Bench 4.0, 13 percentage points above its public-preview result, and 51% on the Artificial Analysis Cyber Index; Mistral also reports a 96.6% average refusal rate on malicious cyber prompts. The model supports three reasoning modes for different task and cost-performance needs; the exact mode names beyond "High" (the default) are not stated.

Availability
The improved model's weights are released; the model had previously been in public preview. Where to obtain the weights, deployment requirements, API availability, pricing, and access regions are not stated.

Why it matters
Teams building coding agents can evaluate a reported 47% reduction in tokens per task against a similar Coding Agent Index score, and a reported 30% reduction in output tokens on Terminal-Bench versus the preview. Open weights also make self-hosting and customization possible, although deployment details are not provided.

Caveats
The benchmark results and efficiency comparisons are reported in Mistral's announcement; independent validation details, benchmark configurations, and confidence intervals are not stated. The article does not specify supported modalities, context length, hardware requirements, weight license, or how to access the weights.

Full Announcement Details

Built for Coders, Trusted by Defenders.
Back to Blog

Daily AI Dev Digest

Saturday, October 31, 2026

Mistral released Mistral Large 4 (Le Chonk), a 1-trillion-parameter multimodal model, now available with improved coding, cybersecurity, and reasoning efficiency. It leads open-weight benchmarks in coding and defensive security while offering three reasoning modes for cost-performance balance. Fully open-weight and European-trained, it empowers sovereign AI with customizable, self-hosted deployment options.

We release the weights of Le Chonk, advancing the frontier of open models in agentic coding and cybersecurity.

Earlier this month, we released Mistral Large 4 (Le Chonk), our flagship 1-trillion-parameter, natively multimodal model, in public preview. The response to the preview has exceeded our expectations, and we're grateful to our developer community, early-access customers, and partners for their support, collaboration, and feedback.

Since then, we've put that feedback to work while continuing to scale reinforcement learning, driving massive improvements in coding and defensive cybersecurity . Gains continue to accelerate, with no signs of saturation.

Today, we're releasing the weights of this improved Mistral Large 4. The new model also features substantial gains in reasoning efficiency that give developers greater performance while lowering inference costs.

Mistral Large 4 ranks as the best open-weight model for coding and among the world's most capable open models overall. Trained from scratch on European infrastructure with clear lineage from inception to release, it gives organizations a trusted foundation for Sovereign AI

Mistral is defining the sovereign AI category and leading the open frontier. We're committed to building the world's most capable models and releasing them openly, giving developers, enterprises, and institutions the power to own, customize, and control their AI.

Built for the Builders.

Mistral Large 4 entered public preview with strong results across leading coding benchmarks. Just three weeks of additional reinforcement learning have since propelled it to the forefront of open-weight coding, taking the lead on FrontierCode and Terminal-Bench while matching the best open-weight models on the Artificial Analysis Coding Agent Index with substantially greater token efficiency.

On Cognition's private FrontierCode 2.0 benchmark, which assesses end-to-end code quality and whether a repository maintainer would merge the agent's change, Mistral Large 4 leads with a score of 19.7%.

Cognition - FrontierCode 2.0

Pushing the pareto frontier for agentic coding.

Mistral Large 4 delivers frontier coding performance with significantly fewer tokens, making sophisticated coding agents more economical to run at scale.

On the Artificial Analysis Coding Agent Index, Mistral Large 4 scores 48%, effectively tied with the leading open-weight model, GLM-5.3, while using 47% fewer tokens per task on average.

Artificial Analysis Coding Agent Index vs. Token Efficiency

Other models

Coding...

Want to learn more?

https://mistral.ai/news/mistral-large-4-oss/

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.