Timestral: Mistral release time-series foundation models

From: Hacker News

Source: Hacker News

The Brief

What shipped
Mistral announced Timestral, its first family of time-series foundation models, in four sizes: 5M, 28M, 123M, and 209M parameters.

Capabilities
Timestral produces probabilistic multivariate forecasts and supports past-only and known-future covariates, including categorical inputs. Mistral says the 209M model ranks first across GIFT-Eval, TIME, and fev-bench, and reports CPU inference on an 8-thread server of 10 ms for the 5M model and 140 ms for the 209M model for a one-day forecast from three weeks of 15-minute history; the 28M model is five times faster than the 209M, while the 123M model is within 0.2% of its score at three-quarters of its latency.

Availability
Not stated whether Timestral is available as downloadable weights, through an API, or by another access route. The announcement mentions domain fine-tuning via Mistral Forge, but does not state availability, pricing, or terms.

Why it matters
The reported CPU timings make it worth benchmarking the smaller checkpoints for high-volume forecasting deployments that want to avoid GPU serving. Covariate support and zero-shot forecasting may also reduce per-domain training work, but validate performance on your own series before replacing a fitted model.

Caveats
The benchmark rankings and speed figures are Mistral-reported; the article says results were submitted to the GIFT-Eval and TIME leaderboards and provides no exact benchmark scores in the supplied text. It does not state inference context limits, weight licensing, deployment requirements beyond the CPU example, or the full training details; the article text also cuts off mid-description of the two-stage training.

Daily AI Dev Digest

Monday, October 19, 2026

Full Announcement Details

Mistral announced Timestral, its first family of time-series foundation models, in four sizes: 5M, 28M, 123M, and 209M parameters. Timestral produces probabilistic multivariate forecasts and supports past-only and known-future covariates, including categorical inputs. Mistral says the 209M model ranks first across GIFT-Eval, TIME, and fev-bench, and reports CPU inference on an 8-thread server of 10 ms for the 5M model and 140 ms for the 209M model for a one-day forecast from three weeks of 15-minute history; the 28M model is five times faster than the 209M, while the 123M model is within 0.2% of its score at three-quarters of its latency.

Want to learn more?

https://mistral.ai/news/timestral/

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.