Together Post-Training: frontier-grade model shaping, the way that fits you

From: Together AI

Source: Together AI

The Brief

What shipped
Together AI introduced Together Post-Training, a platform service for building custom models with supervised fine-tuning (SFT) and reinforcement learning (RL).

Capabilities
Users define the training loop and can run LoRA experiments-including multiple adapters sharing a GPU allocation-or full-weight training for select models. The service supports standard losses including cross-entropy, GRPO, CISPO, and DPPO, custom objectives with gradients sent to its Trainer, and asynchronous RL; launch models include Muse Spark Open, GLM 5.3, Kimi K3, and Qwen 3.8-27B.

Availability
The service is offered on the Together platform, and model resources are billed per GPU-minute; rates are not stated. The article does not state access requirements, regions, or whether availability is preview or general availability.

Why it matters
Teams can bring their own data, environments, and agent harnesses while using Together-managed training and rollout infrastructure, rather than building that stack themselves. Dedicated GPU allocations avoid competing with other users, and the resulting checkpoint can be deployed elsewhere or further refined, according to the article.

Caveats
The article gives no benchmark results, measured speedups, or model-quality comparisons to substantiate its performance claims. Full-weight training is limited to select models, and GPU-minute prices are not provided.

Daily AI Dev Digest

Tuesday, October 20, 2026

Full Announcement Details

Together AI introduced Together Post-Training, a platform service for building custom models with supervised fine-tuning (SFT) and reinforcement learning (RL). Users define the training loop and can run LoRA experiments-including multiple adapters sharing a GPU allocation-or full-weight training for select models. The service supports standard losses including cross-entropy, GRPO, CISPO, and DPPO, custom objectives with gradients sent to its Trainer, and asynchronous RL; launch models include Muse Spark Open, GLM 5.3, Kimi K3, and Qwen 3.8-27B.

Want to learn more?

https://www.together.ai/blog/together-post-training

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.