Daily AI Dev Digest

Wednesday, October 21, 2026

You can now choose 12-hour prompt caching in the API.

From: OpenAI Platform

Source: OpenAI Platform

The Brief

What shipped
OpenAI added selectable 12-hour prompt caching to the API endpoint v1/responses.

Capabilities
The longer cache duration can keep shared instructions, tools, and context cached when requests are hours apart. 12-hour cache writes cost 2.5× the input token price; the announcement says longer caching can benefit lower cached-input pricing.

Availability
Available in the API for v1/responses. The announcement does not state setup instructions, other endpoint support, or further pricing details.

Why it matters
For workloads with requests spaced hours apart, you can choose longer-lived caching for shared prompt content and potentially reduce cached-input costs. Weigh that against the 2.5× input-price charge for cache writes.

Caveats
The announcement does not specify how to enable the duration, cache eligibility or retention behavior beyond 12 hours, or the exact cached-input prices.

Full Announcement Details

OpenAI added selectable 12-hour prompt caching to the API endpoint v1/responses. The longer cache duration can keep shared instructions, tools, and context cached when requests are hours apart. 12-hour cache writes cost 2.5× the input token price; the announcement says longer caching can benefit lower cached-input pricing.

Daily AI Dev Digest

Wednesday, October 21, 2026

Want to learn more?

https://platform.openai.com/docs/changelog

This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.