# Daily AI Dev Digest

*Sunday, October 25, 2026*

# Five Qwen3.8-27B models match Fable 5 Using GVS5H Harness

*From: Hacker News*

Source: Hacker News

---

## The Brief

**What shipped**
The GVS5H authors released a paper and code for ledger-based, zero-shot multi-agent orchestration, evaluating five instances of Qwen3.8-27B on LiveCodeBench Hard.

**Capabilities**
On the 100 latest hard LiveCodeBench problems, the reported Qwen3.8-27B result rises from 66.8% pass@1 alone to 92.4% with orchestration, versus 90.4% for Fable 5. The paper also reports Qwen3.8 Flash Next at 93.0% pass@1 and 9% of Fable 5's cost; that cost claim is for Flash Next, not the 27B model.

**Availability**
The repository includes the benchmark runner and is MIT-licensed; running it requires uv and an API key for the chosen model, with provider options listed in the instructions. Model-weight download details, concrete prices, and deployment requirements are not stated.

**Why it matters**
For coding-agent evaluations, the harness is a way to test whether multiple fresh model instances coordinating through shared files improve pass@1 over a single call. The reported 27B result makes local serving worth evaluating if your workload can tolerate the orchestration and its additional inference, but its cost and latency are not stated.

**Caveats**
The results cover 100 benchmark problems, and the authors say gains are not universal. Opus 5 was tested on an earlier scaffold in a single OpenRouter pass rather than with five agents on a pinned backend, so that comparison is not like-for-like; the reported results are not independently verified here.

---

# Daily AI Dev Digest

*Sunday, October 25, 2026*

### Full Announcement Details

*The GVS5H authors released a paper and code for ledger-based, zero-shot multi-agent orchestration, evaluating five instances of Qwen3.8-27B on LiveCodeBench Hard. On the 100 latest hard LiveCodeBench problems, the reported Qwen3.8-27B result rises from 66.8% pass@1 alone to 92.4% with orchestration, versus 90.4% for Fable 5. The paper also reports Qwen3.8 Flash Next at 93.0% pass@1 and 9% of Fable 5's cost; that cost claim is for Flash Next, not the 27B model.*

### Want to learn more?

https://github.com/slee-persis/GVS5H

*This digest tracks the latest developments from the frontier AI labs - model releases, product and API launches, research reports, and capability results - as they ship.*